

# ECON 425 HW9

March 5, 2024

## Problem 2 (K-means clustering)

In this problem, you will perform K-means clustering manually, with  $K = 2$ , on a small example with  $n = 6$  observations and  $p = 2$  features. The data matrix is

$$X = \begin{pmatrix} 1 & 4 \\ 1 & 3 \\ 0 & 4 \\ 5 & 1 \\ 6 & 2 \\ 4 & 0 \end{pmatrix}$$

- (a) Plot the observations.
- (b) Randomly assign a cluster label to each observation. You can use the `np.random.choice()` function to do this. Report the cluster labels for each observation.
- (c) Compute the centroid for each cluster.
- (d) Assign each observation to the centroid to which it is closest, in terms of Euclidean distance. Report the cluster labels for each observation.
- (e) Repeat (c) and (d) until the answers obtained stop changing.
- (f) In your plot from (a), color the observations according to the cluster labels obtained.

## Problem 3 (PCA on Olivetti faces)

Use the Olivetti faces dataset available through sklearn to do the following.

- (a) Fetch and load the data with the `fetch_olivetti_faces` method from `sklearn.datasets`.
- (b) Demean each face in the data set (no need to divide by standard deviation as every dimension is a number between a fixed range representing a pixel).
- (c) Compute and display the first 9 *eigenfaces*. The  $k$ -th eigenface of a given face is an image based on the first  $k$  principal components only.
- (d) Any given face in the data set can be represented as a linear combination of the eigenfaces. For any face in the data set, show how it progresses as we combine 1, 51, 101, ... eigenfaces, until the full image is recovered.

# ML HW 9

March 19, 2024

```
[40]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```


## 0.1 Problem 2

```
[41]: df = pd.DataFrame({"x1": [1,1,0,5,6,4], "x2": [4,3,4,1,2,0]})
```

### 0.1.1 a)

Below I plot the observations from the dataframe.

```
[42]: df.plot(x= "x1", y="x2", kind='scatter')
plt.show()
```



### 0.1.2 b)

```
[46]: for i in df.index:
        df.loc[i, 'cluster'] = np.random.choice([1,2])
        print("Observation " + str(i+1) + " is assigned to cluster " + str(int(df.
        ↪loc[i, 'cluster'])))
```

Observation 1 is assigned to cluster 2  
Observation 2 is assigned to cluster 2  
Observation 3 is assigned to cluster 1  
Observation 4 is assigned to cluster 1  
Observation 5 is assigned to cluster 2  
Observation 6 is assigned to cluster 1

### 0.1.3 c)

```
[47]: c1x1= df[df.cluster==1]["x1"].mean()
c1x2= df[df.cluster==1]["x2"].mean()
c2x1= df[df.cluster==2]["x1"].mean()
c2x2= df[df.cluster==2]["x2"].mean()

print(f"The x-coordinate for centroid 1 is {c1x1} and the y-coordinate for_
↪centroid 1 is {c1x2}")
print(f"The x-coordinate for centroid 2 is {c2x1} and the y-coordinate for_
↪centroid 2 is {c2x2}")
```

The x-coordinate for centroid 1 is 3.0 and the y-coordinate for centroid 1 is 1.6666666666666667

The x-coordinate for centroid 2 is 2.6666666666666665 and the y-coordinate for centroid 2 is 3.0

### 0.1.4 d)

Below I assign each observation to the centroid to which it is closest, in terms of Euclidean distance.

```
[48]: for i in df.index:
        dc1= np.sqrt((df.loc[i, "x1"]-c1x1)**2 + (df.loc[i, "x2"]-c1x2)**2)
        dc2= np.sqrt((df.loc[i, "x1"]-c2x1)**2 + (df.loc[i, "x2"]-c2x2)**2)

        if dc1<dc2:
            df.loc[i, "cluster"] =1
            print(f"Observation {i+1} is now in cluster 1")
        else:
            df.loc[i, "cluster"] =2
            print(f"Observation {i+1} is now in cluster 2")
```

Observation 1 is now in cluster 2  
 Observation 2 is now in cluster 2  
 Observation 3 is now in cluster 2  
 Observation 4 is now in cluster 1  
 Observation 5 is now in cluster 1  
 Observation 6 is now in cluster 1

### 0.1.5 e)

Below I repeat c) and d) until the answers stop changing. We can observe that after one iteration in the while loop, the cluster assignments are still the same as in d).

```
[49]: check = [0,0,0,0,0,0]

while (check != df.cluster).all():

    check = df['cluster'].copy()

    c1x1= df[df.cluster==1]["x1"].mean()
    c1x2= df[df.cluster==1]["x2"].mean()
    c2x1= df[df.cluster==2]["x1"].mean()
    c2x2= df[df.cluster==2]["x2"].mean()

    for i in df.index:
        dc1= np.sqrt((df.loc[i,"x1"]-c1x1)**2 + (df.loc[i,"x2"]-c1x2)**2)
        dc2= np.sqrt((df.loc[i,"x1"]-c2x1)**2 + (df.loc[i,"x2"]-c2x2)**2)

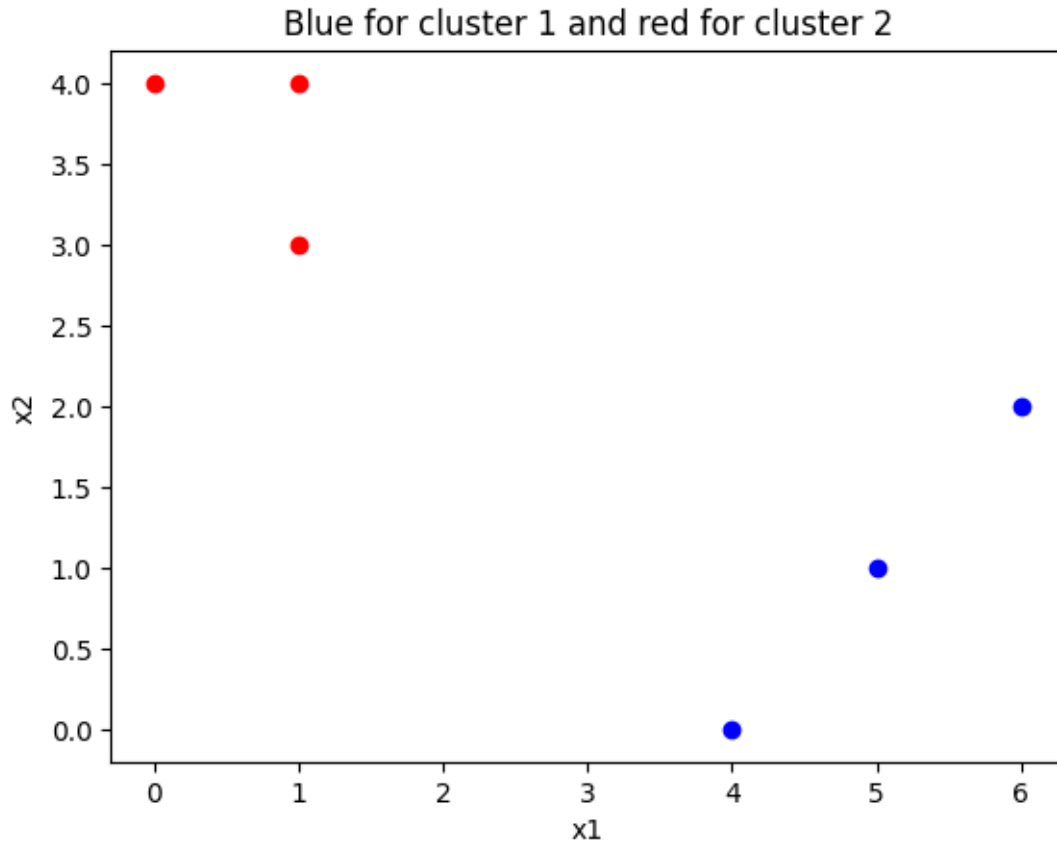
        if dc1<dc2:
            df.loc[i,"cluster"] =1
            print(f"Observation {i+1} is now in cluster 1")
        else:
            df.loc[i,"cluster"] =2
            print(f"Observation {i+1} is now in cluster 2")
```

Observation 1 is now in cluster 2  
 Observation 2 is now in cluster 2  
 Observation 3 is now in cluster 2  
 Observation 4 is now in cluster 1  
 Observation 5 is now in cluster 1  
 Observation 6 is now in cluster 1

### 0.1.6 f)

```
[53]: c1 =df[df.cluster==1]
      c2 =df[df.cluster==2]
      plt.scatter(c1.x1, c1.x2, color='blue')
      plt.scatter(c2.x1,c2.x2, color='red')
```

```
plt.xlabel("x1")
plt.ylabel("x2")
plt.title("Blue for cluster 1 and red for cluster 2")
plt.show()
```



## 0.2 Problem 3

### 0.2.1 a)

Below I fetch and load the data with the `fetch_olivetti_faces` method from `sklearn.datasets`

```
[66]: from sklearn.datasets import fetch_olivetti_faces

faces = fetch_olivetti_faces()

df = pd.DataFrame(faces.data)
```

### 0.2.2 b)

I then demean each face in the dataset.

```
[67]: demean = df.mean(axis=0)
```

```
[68]: df = df.sub(demean, axis = 1)
```

### 0.2.3 c)


Below I will compute 9 eigenfaces using PCA.



```
[85]: from sklearn.decomposition import PCA
```


```
ef = 9  
pca = PCA(n_components=ef)  
pca.fit(df)
```

```
[85]: PCA(n_components=9)
```

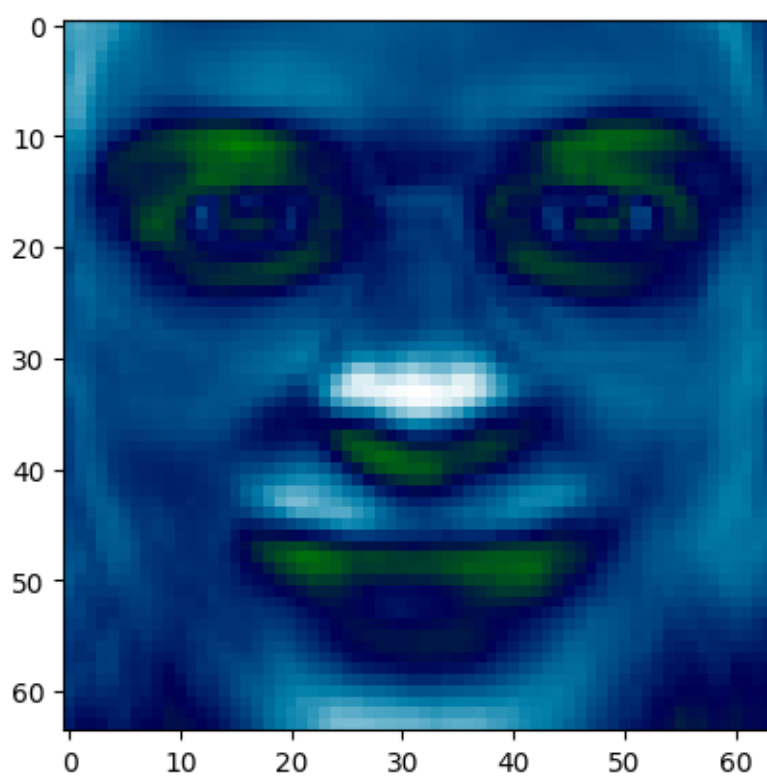

```
[127]: for i in range(9):  
        plt.imshow(pca.components_[i].reshape(64, 64), cmap='ocean')  
        plt.show()
```


















### 0.2.4 d)

Below I reconstruct the first face. I begin by storing the first face in face1.


```
[104]: face1= df.iloc[0]
```

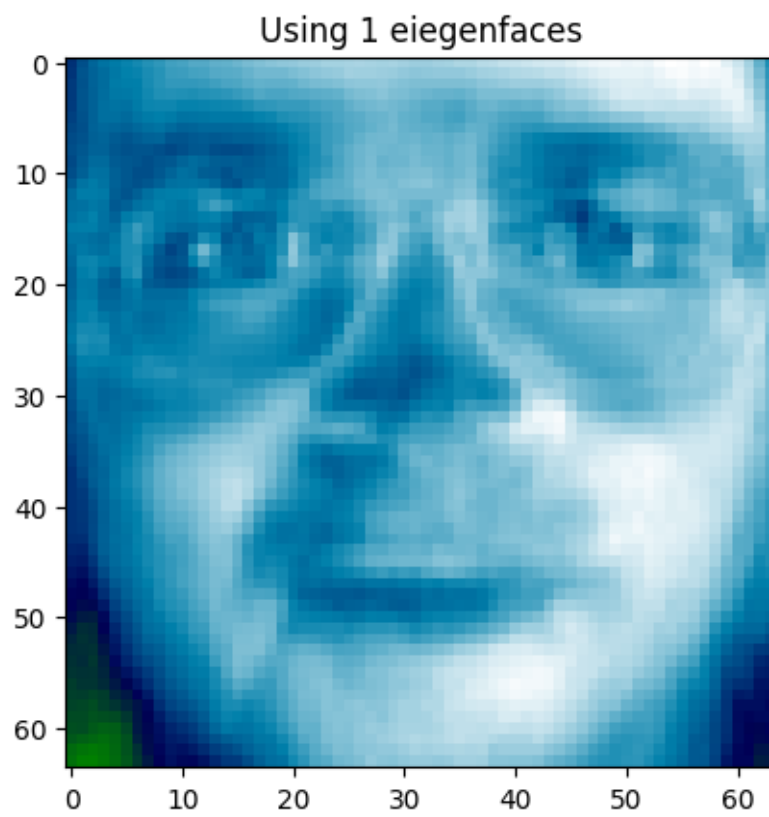
Then, for each number of eigenfaces, i fit a new PCA model and plot the result.


```
[124]: efs = [1, 51, 101, 151, 201, 251, 301, 351]

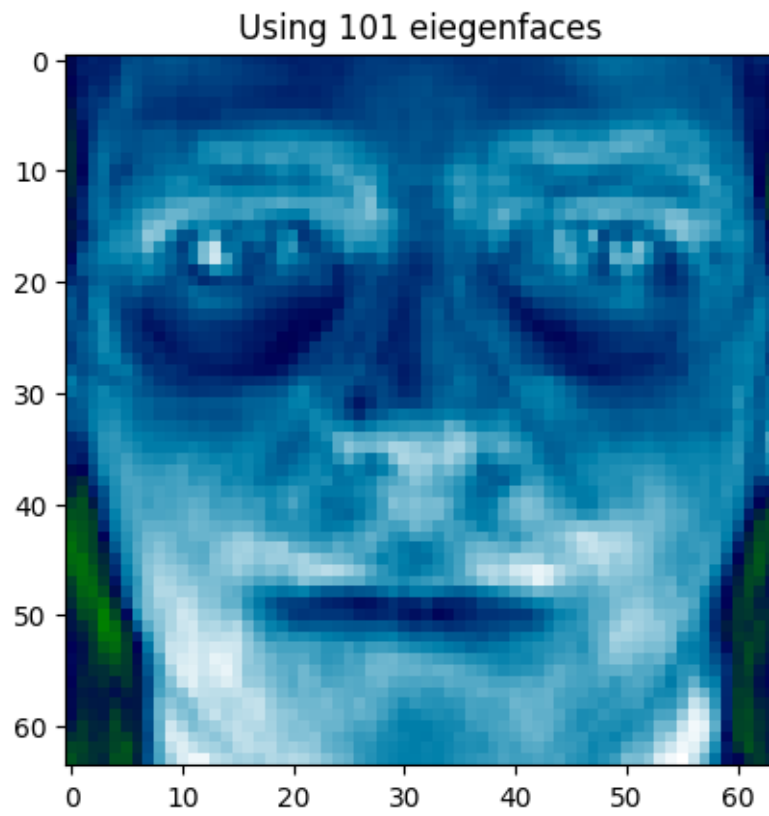
plt.imshow(face1.values.reshape(64, 64), cmap='ocean')
plt.title("Original")
plt.show()


for comps in efs:
    cur_pca = PCA(n_components=comps)
    cur_pca.fit(df)
    cur_comps = cur_pca.transform(face1.values.reshape(1, -1))
    cur_face = cur_pca.inverse_transform(cur_comps)
    plt.imshow(cur_face.reshape(64, 64), cmap='ocean')
    plt.title(f"Using {comps} eigenfaces")
    plt.show()
```



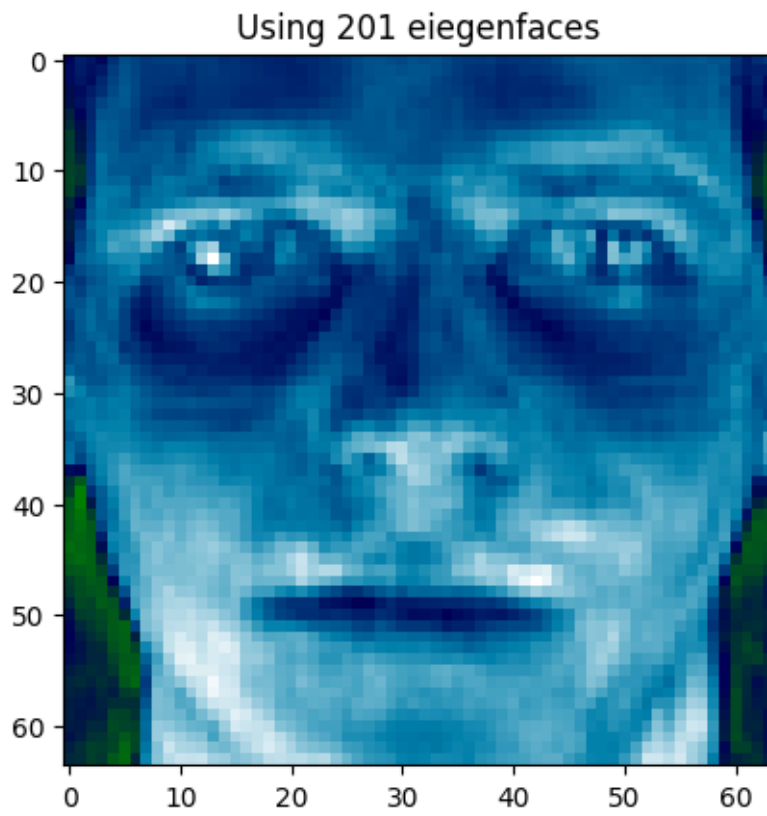





















[ ]: