

**Problem 2 ( fitting a simple neural network)**

Use the dataset card transdata.csv from the previous homework and maintain the same train-test split.

Fit a feedforward neural network with two ReLU layers using stochastic gradient descent (SGD). Follow [this tutorial](#). Experiment with the number of neurons per layer, the number of epochs, the learning rate for SGD, and the batch size for backpropagation. Report accuracy and F1 score on the test sample. Does your model perform better than a simple decision tree from the last homework?

**0.2 Problem 2**

Below I download the card dataset and split it into 50% for training and 50% for testing.

```
[11]: df =pd.read_csv("/Users/sandinatatu/Desktop/card_transdata-1-2.csv")

X=df.drop("fraud", axis=1).values
y=df["fraud"].values

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = .5,
↳random_state = 10)
```

I then transform my training and testing datasets into torch object. After doing so, I create a function detailing the structure of my neural network (the number of neurons per hidden layer will be detailed by a variable I called “hidden\_neurons”).

```
[12]: import torch
from torch import nn
from torch.utils.data import DataLoader
from torchvision import datasets
from torchvision.transforms import ToTensor
from torch.utils.data import DataLoader, TensorDataset

X_train_tensor = torch.tensor(X_train, dtype=torch.float32)
y_train_tensor = torch.tensor(y_train, dtype=torch.float32)
X_test_tensor = torch.tensor(X_test, dtype=torch.float32)
y_test_tensor = torch.tensor(y_test, dtype=torch.float32)

train_dataset = TensorDataset(X_train_tensor, y_train_tensor)
train_dataloader = DataLoader(train_dataset, batch_size = 64)
```

```

test_dataset = TensorDataset(X_test_tensor, y_test_tensor)
test_dataloader = DataLoader(test_dataset, batch_size=64)

class NeuralNetwork(nn.Module):
    def __init__(self, hn):
        super().__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(7, hn),
            nn.ReLU(),
            nn.Linear(hn, hn),
            nn.ReLU(),
            nn.Linear(hn, 2),
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits

```

I then use the `train_loop` and the `test_loop` from the PyTorch website. These loops will help with understanding the performance on the training data, as well as on the test data.

```

[13]: from sklearn.metrics import f1_score

def train_loop(dataloader, model, loss_fn, optimizer):
    size = len(dataloader.dataset)

    model.train()
    for batch, (X, y) in enumerate(dataloader):

        y=y.long()
        pred = model(X)
        loss = loss_fn(pred, y)

        loss.backward()
        optimizer.step()
        optimizer.zero_grad()

        if batch % 100 == 0:
            loss, current = loss.item(), batch * batch_size + len(X)

def test_loop(dataloader, model, loss_fn):

```

```

model.eval()
size = len(dataloader.dataset)
num_batches = len(dataloader)
test_loss, correct = 0, 0

f1s=[]

with torch.no_grad():
    for X, y in dataloader:
        y=y.long()
        pred = model(X)
        test_loss += loss_fn(pred, y).item()
        correct += (pred.argmax(1) == y).type(torch.float).sum().item()
        cf1=f1_score(y.cpu(), pred.argmax(1).cpu(), average='weighted')
        f1s.append(cf1)

af1= sum(f1s) / num_batches
test_loss /= num_batches
correct /= size
print(f"Test Error: \n Accuracy: {(100*correct):>0.1f}%, F1 Score: {af1}, \n
↪Avg loss: {test_loss:>8f} \n")

```

### 0.2.1 Model 1

For the first model, I will use 100 neurons per hidden layer. the learning rate will be 1e-3, the batch size will be 64, and I will have 5 epochs.

```

[14]: learning_rate = 1e-3
      batch_size = 64
      epochs = 5
      hn = 100

      model = NeuralNetwork(hn)

```

```

[15]: loss_fn = nn.CrossEntropyLoss()

```

```

[16]: optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)

```

```

[17]: loss_fn = nn.CrossEntropyLoss()
      optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)

      for t in range(epochs):
          print(f"Epoch {t+1}\n-----")

```

```
train_loop(train_dataloader, model, loss_fn, optimizer)
test_loop(test_dataloader, model, loss_fn)
print("Done!")
```

Epoch 1

-----  
Test Error:

Accuracy: 92.3%, F1 Score: 0.9160726763778203, Avg loss: 0.223735

Epoch 2

-----  
Test Error:

Accuracy: 92.9%, F1 Score: 0.923822153040678, Avg loss: 0.182242

Epoch 3

-----  
Test Error:

Accuracy: 93.5%, F1 Score: 0.9288238461945655, Avg loss: 0.160974

Epoch 4

-----  
Test Error:

Accuracy: 94.3%, F1 Score: 0.9385039195805155, Avg loss: 0.138886

Epoch 5

-----  
Test Error:

Accuracy: 95.3%, F1 Score: 0.950725192053029, Avg loss: 0.121787

Done!

### 0.2.2 Model 2

For the second model, I will use 50 neurons per hidden\_layer, a batch\_size of 100, and 50 epochs. I will leave the learning rate the same.

```
[18]: learning_rate = 1e-3
      batch_size = 100
      epochs = 50
      hn = 50

      model = NeuralNetwork(hn)
```

```
[19]: loss_fn = nn.CrossEntropyLoss()
```

```
[20]: optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)
```

```
[21]: loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)

for t in range(epochs):
    print(f"Epoch {t+1}\n-----")
    train_loop(train_dataloader, model, loss_fn, optimizer)
    test_loop(test_dataloader, model, loss_fn)
print("Done!")
```

Epoch 1

-----

Test Error:

Accuracy: 92.3%, F1 Score: 0.9132063914546926, Avg loss: 0.200240

Epoch 2

-----

Test Error:

Accuracy: 93.0%, F1 Score: 0.9226235391808396, Avg loss: 0.165299

Epoch 3

-----

Test Error:

Accuracy: 93.5%, F1 Score: 0.9284519850762067, Avg loss: 0.148936

Epoch 4

-----

Test Error:

Accuracy: 94.4%, F1 Score: 0.9390439548902482, Avg loss: 0.130812

Epoch 5

-----

Test Error:

Accuracy: 95.2%, F1 Score: 0.949355832975581, Avg loss: 0.116275

Epoch 6

-----

Test Error:

Accuracy: 95.6%, F1 Score: 0.9544338363229107, Avg loss: 0.104019

Epoch 7

-----

Test Error:

Accuracy: 96.1%, F1 Score: 0.9591953420480193, Avg loss: 0.093551

Epoch 8

-----

Test Error:

Accuracy: 96.4%, F1 Score: 0.9634182365142981, Avg loss: 0.085939

Epoch 9

-----  
Test Error:

Accuracy: 96.6%, F1 Score: 0.9652521850493413, Avg loss: 0.079909

Epoch 10

-----  
Test Error:

Accuracy: 96.8%, F1 Score: 0.9678637326844176, Avg loss: 0.076111

Epoch 11

-----  
Test Error:

Accuracy: 97.0%, F1 Score: 0.9706895906526415, Avg loss: 0.073264

Epoch 12

-----  
Test Error:

Accuracy: 97.2%, F1 Score: 0.9725246888522892, Avg loss: 0.070227

Epoch 13

-----  
Test Error:

Accuracy: 97.4%, F1 Score: 0.9748575221745271, Avg loss: 0.064727

Epoch 14

-----  
Test Error:

Accuracy: 97.5%, F1 Score: 0.9751878905518522, Avg loss: 0.063259

Epoch 15

-----  
Test Error:

Accuracy: 97.8%, F1 Score: 0.9787445076346779, Avg loss: 0.059120

Epoch 16

-----  
Test Error:

Accuracy: 97.9%, F1 Score: 0.9791231415936769, Avg loss: 0.057238

Epoch 17

-----  
Test Error:

Accuracy: 97.9%, F1 Score: 0.9793056718451277, Avg loss: 0.056256

Epoch 18

Test Error:

Accuracy: 97.9%, F1 Score: 0.9795504174867682, Avg loss: 0.056023

Epoch 19

-----

Test Error:

Accuracy: 98.2%, F1 Score: 0.982315364964906, Avg loss: 0.051916

Epoch 20

-----

Test Error:

Accuracy: 98.2%, F1 Score: 0.982187012946666, Avg loss: 0.053296

Epoch 21

-----

Test Error:

Accuracy: 98.0%, F1 Score: 0.9805450132548946, Avg loss: 0.050344

Epoch 22

-----

Test Error:

Accuracy: 98.3%, F1 Score: 0.9831015023112061, Avg loss: 0.045974

Epoch 23

-----

Test Error:

Accuracy: 98.5%, F1 Score: 0.9851329637338113, Avg loss: 0.044682

Epoch 24

-----

Test Error:

Accuracy: 98.3%, F1 Score: 0.983126386841809, Avg loss: 0.045271

Epoch 25

-----

Test Error:

Accuracy: 98.5%, F1 Score: 0.9853138899429756, Avg loss: 0.042274

Epoch 26

-----

Test Error:

Accuracy: 98.4%, F1 Score: 0.9843099309057519, Avg loss: 0.042933

Epoch 27

-----

Test Error:

Accuracy: 98.4%, F1 Score: 0.9847836854087625, Avg loss: 0.044477

Epoch 28

-----

Test Error:

Accuracy: 98.5%, F1 Score: 0.9859161609136025, Avg loss: 0.041434

Epoch 29

-----

Test Error:

Accuracy: 98.5%, F1 Score: 0.9855231985343685, Avg loss: 0.040554

Epoch 30

-----

Test Error:

Accuracy: 98.4%, F1 Score: 0.9849782514459564, Avg loss: 0.040986

Epoch 31

-----

Test Error:

Accuracy: 98.9%, F1 Score: 0.9890392173414442, Avg loss: 0.040421

Epoch 32

-----

Test Error:

Accuracy: 99.0%, F1 Score: 0.9894246893447735, Avg loss: 0.035417

Epoch 33

-----

Test Error:

Accuracy: 98.9%, F1 Score: 0.989154000380651, Avg loss: 0.035672

Epoch 34

-----

Test Error:

Accuracy: 98.9%, F1 Score: 0.9897297163755264, Avg loss: 0.035281

Epoch 35

-----

Test Error:

Accuracy: 98.9%, F1 Score: 0.98969990894516, Avg loss: 0.034250

Epoch 36

-----

Test Error:

Accuracy: 98.9%, F1 Score: 0.9896103473092636, Avg loss: 0.033170

Epoch 37

-----

Test Error:



Accuracy: 99.1%, F1 Score: 0.99095657155999, Avg loss: 0.034581

Epoch 38

-----  
Test Error:

Accuracy: 99.3%, F1 Score: 0.9927376872614276, Avg loss: 0.030884

Epoch 39

-----  
Test Error:

Accuracy: 99.0%, F1 Score: 0.9901989634513203, Avg loss: 0.032929

Epoch 40

-----  
Test Error:

Accuracy: 98.9%, F1 Score: 0.9894375101485104, Avg loss: 0.033249

Epoch 41

-----  
Test Error:

Accuracy: 99.1%, F1 Score: 0.9913511412120248, Avg loss: 0.030650

Epoch 42

-----  
Test Error:

Accuracy: 99.3%, F1 Score: 0.9925677619474407, Avg loss: 0.027472

Epoch 43

-----  
Test Error:

Accuracy: 99.3%, F1 Score: 0.9925659823947043, Avg loss: 0.026487

Epoch 44

-----  
Test Error:

Accuracy: 99.3%, F1 Score: 0.9932099928835201, Avg loss: 0.026834

Epoch 45

-----  
Test Error:

Accuracy: 99.4%, F1 Score: 0.994034604599448, Avg loss: 0.025856

Epoch 46

-----  
Test Error:

Accuracy: 99.4%, F1 Score: 0.9934338583376484, Avg loss: 0.024943

Epoch 47

-----  
Test Error:

Accuracy: 99.5%, F1 Score: 0.994762596644231, Avg loss: 0.024202

Epoch 48

-----  
Test Error:

Accuracy: 99.4%, F1 Score: 0.9934740447645186, Avg loss: 0.024598

Epoch 49

-----  
Test Error:

Accuracy: 99.3%, F1 Score: 0.9934613049869524, Avg loss: 0.026757

Epoch 50

-----  
Test Error:

Accuracy: 99.4%, F1 Score: 0.994350325256214, Avg loss: 0.023354

Done!

### 0.2.3 Model 3

For my final model, I will use 100 neurons per hidden layer, a batch\_size of 64, 10 epochs, and a learning rate of 1e-4.

```
[22]: learning_rate = 1e-4
      batch_size = 64
      epochs = 10
      hn = 100

      model = NeuralNetwork(hn)
```

```
[23]: loss_fn = nn.CrossEntropyLoss()
```

```
[24]: optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)
```

```
[25]: loss_fn = nn.CrossEntropyLoss()
      optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)

      for t in range(epochs):
          print(f"Epoch {t+1}\n-----")
          train_loop(train_dataloader, model, loss_fn, optimizer)
          test_loop(test_dataloader, model, loss_fn)
      print("Done!")
```

Epoch 1

-----  
Test Error:

Accuracy: 91.5%, F1 Score: 0.882261473644765, Avg loss: 0.387805

Epoch 2

-----

Test Error:

Accuracy: 91.7%, F1 Score: 0.8921798371616224, Avg loss: 0.292035

Epoch 3

-----

Test Error:

Accuracy: 91.9%, F1 Score: 0.8984379999682679, Avg loss: 0.243960

Epoch 4

-----

Test Error:

Accuracy: 92.0%, F1 Score: 0.9029252995191686, Avg loss: 0.221302

Epoch 5

-----

Test Error:

Accuracy: 92.2%, F1 Score: 0.9064863137793205, Avg loss: 0.208320

Epoch 6

-----

Test Error:

Accuracy: 92.3%, F1 Score: 0.908930732209424, Avg loss: 0.200694

Epoch 7

-----

Test Error:

Accuracy: 92.4%, F1 Score: 0.9106543366623544, Avg loss: 0.194982

Epoch 8

-----

Test Error:

Accuracy: 92.5%, F1 Score: 0.9121233702504971, Avg loss: 0.188377

Epoch 9

-----

Test Error:

Accuracy: 92.6%, F1 Score: 0.9138915418022129, Avg loss: 0.184001

Epoch 10

-----

Test Error:

Accuracy: 92.7%, F1 Score: 0.9154081861548743, Avg loss: 0.180168

Done!

The highest F1 score is obtained for the second model and is around 99.4. The highest accuracy is also obtained for the second model and is around 99.4. None of the models perform better than the decision tree from last week's homeworks.